

Pictures

Arthur Charpentier

UQAM

Actuarial Summer School 2019

Statistics & ML Toolbox : Neural Nets

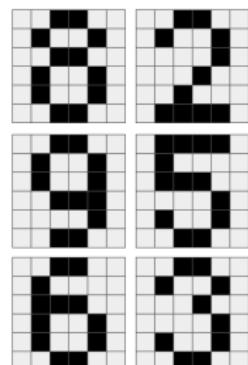
Very popular for pictures...

Picture x_i is

- a $n \times n$ matrix in $\{0, 1\}^{n^2}$ for black & white
- a $n \times n$ matrix in $[0, 1]^{n^2}$ for grey-scale
- a $3 \times n \times n$ array in $([0, 1]^3)^{n^2}$ for color
- a $T \times 3 \times n \times n$ tensor in $(([0, 1]^3)^T)^{n^2}$ for video

y here is the label ("8", "9", "6", etc)

Suppose we want to recognize a "6" on a picture



$$m(x) = \begin{cases} +1 & \text{if } x \text{ is a "6"} \\ -1 & \text{otherwise} \end{cases}$$

Statistics & ML Toolbox : Neural Nets

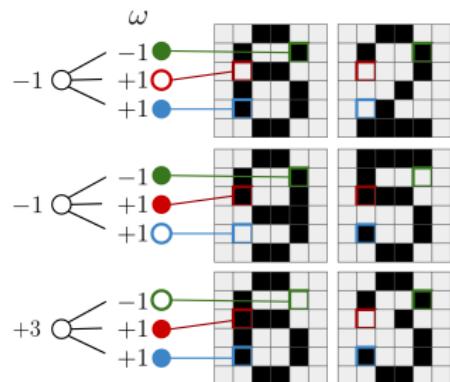
Consider some specific pixels, and associate weights ω such that

$$\hat{m}(\mathbf{x}) = \text{sign} \left(\sum_{i,j} \omega_{i,j} x_{i,j} \right)$$

where

$$x_{i,j} = \begin{cases} +1 & \text{if pixel } x_{i,j} \text{ is black} \quad \blacksquare \\ -1 & \text{if pixel } x_{i,j} \text{ is white} \quad \square \end{cases}$$

for some weights $\omega_{i,j}$ (that can be negative...)

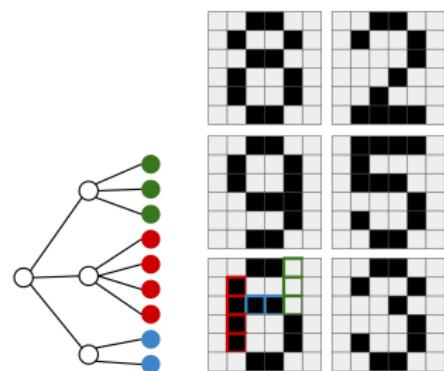


Statistics & ML Toolbox : Neural Nets

A deep network is a network with a lot of layers

$$\hat{m}(\mathbf{x}) = \text{sign} \left(\sum_i \omega_i \hat{m}_i(\mathbf{x}) \right)$$

where \hat{m}_i 's are outputs of previous neural nets. Those layers can capture shapes in some areas nonlinearities, cross-dependence, etc



Detection & Recognition



can be complicated... see Marr (1982, [Vision](#))

Convolutional Network

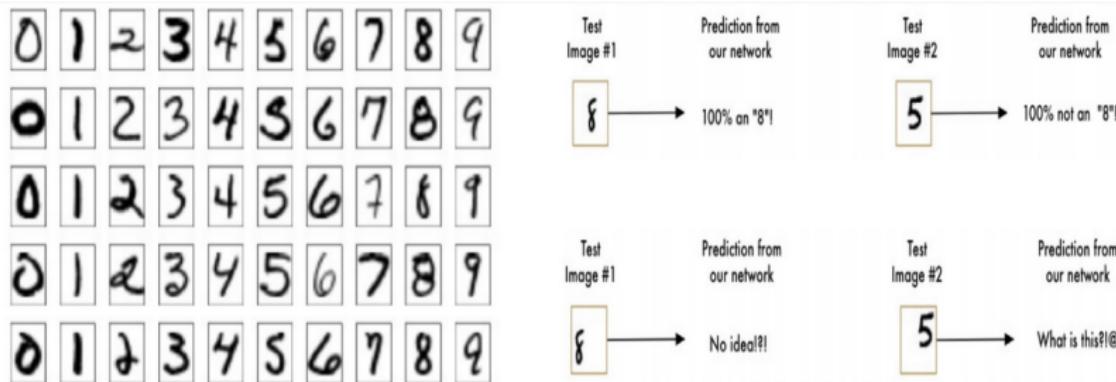


FIGURE 11.9. Examples of training cases from ZIP code data. Each image is a 16×16 8-bit grayscale representation of a handwritten digit.

source: Hastie *et al.* (2009, [The Elements of Statistical Learning](#))

Neural Nets fails to recognize '8' when the digit is not centered
translation, scale and (small) rotation invariances are needed

Rotation, Scaling, Mixing, etc

See <http://yann.lecun.com/exdb/lenet/>

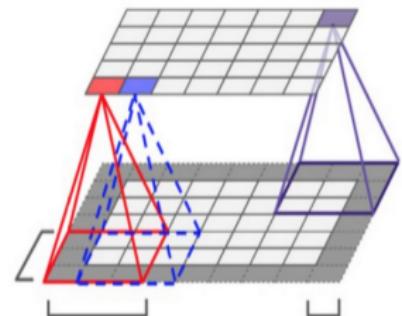
See also <http://scs.ryerson.ca/~aharley/vis/conv>

Convolutional Network

Break the image into overlapping image tiles and, feed each image tile into a small neural network with the same weights (and the same activation function)



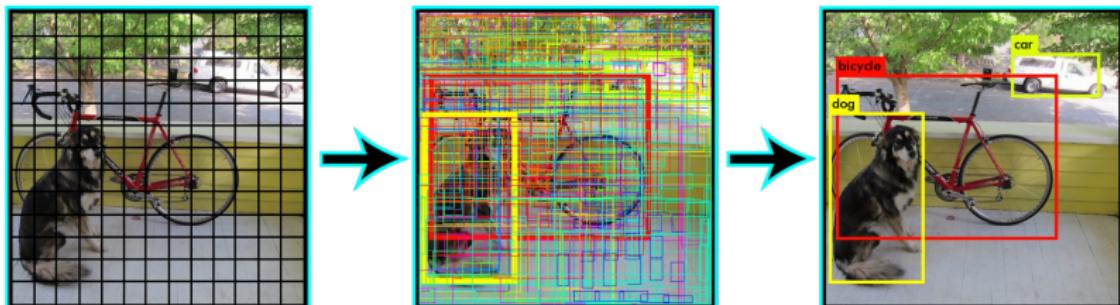
1. ConvNets exploit spatially local correlation: each neuron is locally-connected (to only a small region of the input volume)
2. Reduce the size of the array, using a pooling algorithm.
(pooling step reduces the dimensionality of each feature)



Convolutional Network

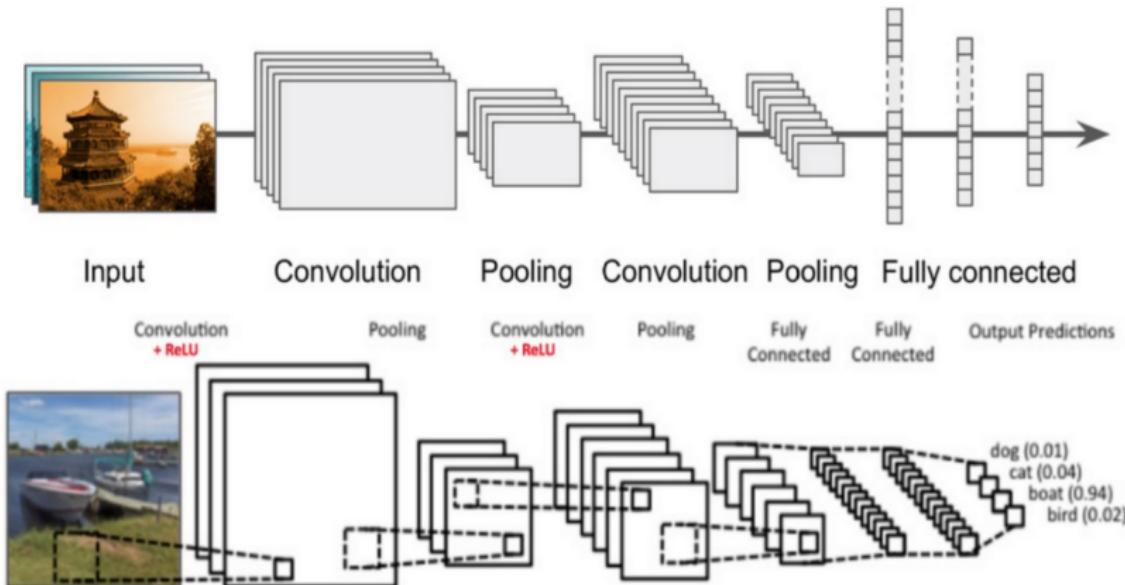
pooling:

- makes the input representations smaller and more manageable
- reduces the number of weights and links in the network, therefore, controlling overfitting
- makes the network invariant to small transformations, distortions and translations in the input image
- helps us arrive at an almost scale invariant representation of our image



(via [yolo](#) for objects detection)

Convolutional Network



Convolutional Network

Can be used to detect faces on a picture



via LeCun (2016, [Leçon inaugurale](#))

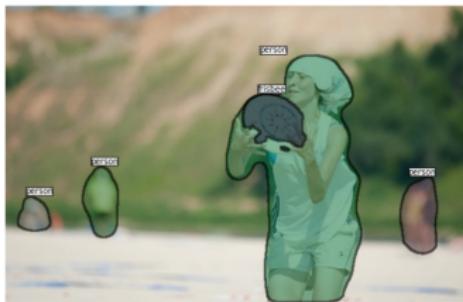
Convolutional Network

and for object detection (even on a video)



via LeCun (2016, [Leçon inaugurale](#))

Convolutional Network



via LeCun (2016, [Leçon inaugurale](#))

Convolutional Network

In some application, need to find scale invariant transformations



via LeCun (2016, [Leçon inaugurale](#))

Recognition with Neural Nets

Mid-2000s [caltech101](#) dataset, with 101 categories, 30 pictures per category (training sample)

2012 [ImageNet](#) dataset, 1.2 million (training) pictures, 20,000 categories see [Li Fei-Fei's page](#)

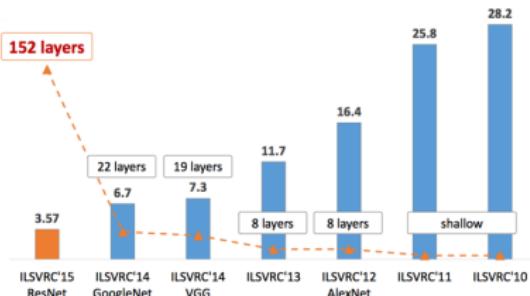
[OverFeat](#) 2012 13.8% error,

[GoogLeNet](#) 2014 6.6%

(22 layers),

[ResNet](#) 2015 3.6% (152 layers)

see [@sidereal's post](#)



From Pictures to Videos

A picture is a $w \times h \times 3$ tensor

A video is a $w \times h \times 3 \times t$ tensor

One can use RNN

recurrent neural nets (or LSTM)

Can also be used to identify objects

see <https://towardsdatascience.com>

Which Pictures? Satellite



From <https://unequalscenes.com/>, Hout Bay, about 15km south of Cape Town, South Africa and Dar es Salaam, Tanzania.

Which Pictures? Satellite



From <https://nytimes.com/>, Chennai (Madras) in India, 2018 and 2019

Which Pictures? Satellite

via <https://geoweb.iau-idf.fr/> (1949 on the left, 2019 on the right)



Which Pictures? Satellite

Possible use for crop insurance:

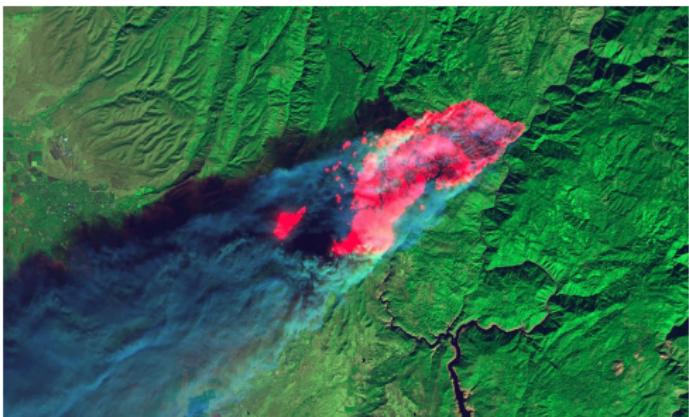
- To assess crop share (relative proportions) in a large area (no georeference available of the fields)
- To estimate yield of an specific crop/season in a large area (no georeference available of the fields)
- To detect and to georeference fields with specific crops
- To detect kind-of-crop info from specific fields

Various sources can be used :

- **MODIS** (Moderate Resolution Imaging Spectroradiometer)
250m × 250m, 2 images per day
- **LANDSAT**, 15m × 15m, 1 image every 16 days
(see also <https://landsat.visibleearth.nasa.gov>)

Which Pictures? Satellite

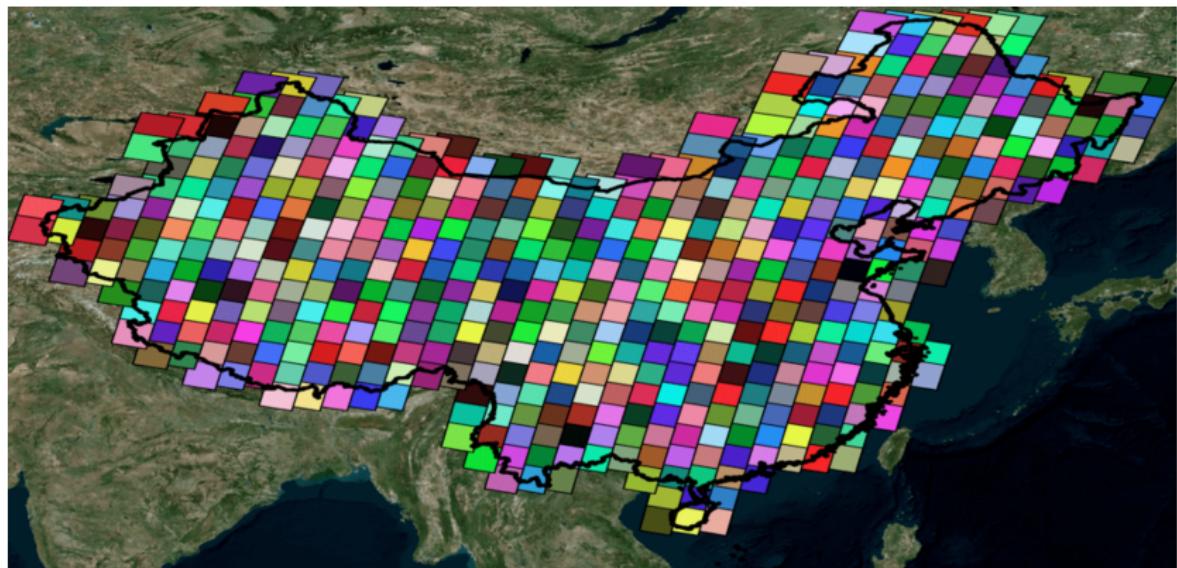
Moderate Resolution Imaging Spectroradiometer (MODIS) on NASA's Aqua satellite of the Camp Fire on November 12, 2018.
[Photo: NASA Earth Observatory vs Nov. 8, 2018, by [Landsat 8](#), shows short-wave infrared (red), which gives the full extent of the actively burning area of the Camp Fire]



Which Pictures? Satellite

Landsat picture = 17,500px × 14,500, 950Mb

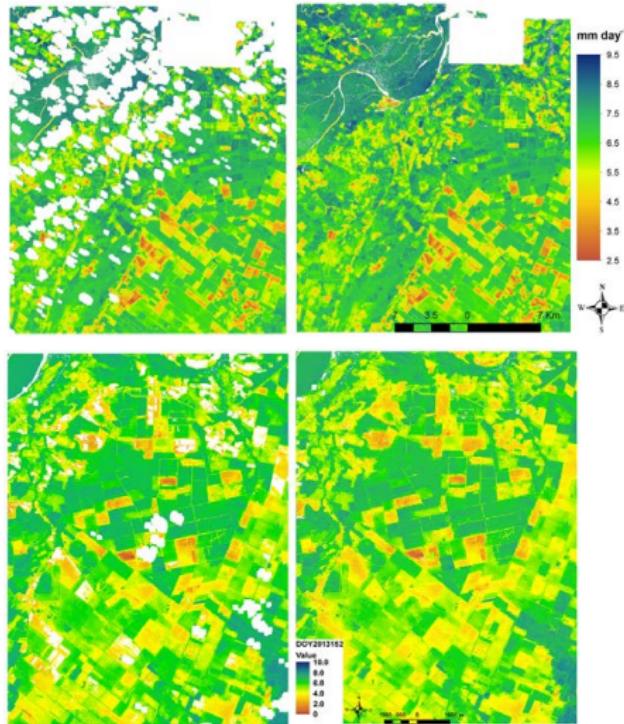
Landsat 8 to cover China : 530 tiles, 23 per year \sim 11Tb per year



Which Pictures? Satellite

Satellite pictures are
more than visible colors
TIRS (Thermal Infrared Sensor)

Practical problem : clouds
(see <https://forestthreats.org/>
or <https://harrisgeospatial.com/>)

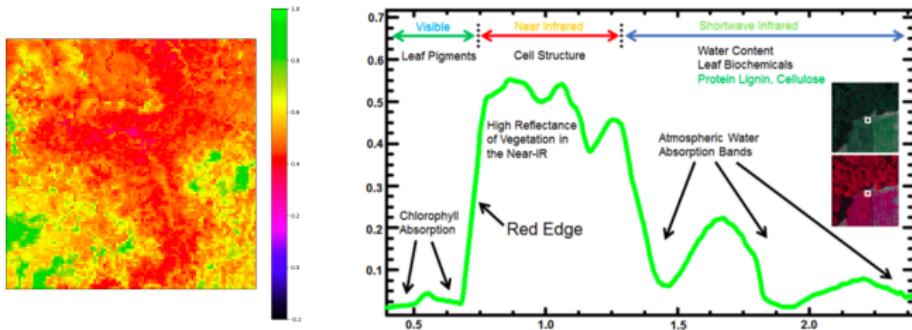


Which Pictures? Satellite

Elowitz (2013, [What is Imaging Spectroscopy \(Hyperspectral Imaging\)?](#))

Infrareds used to produce (e.g.) some Normalized Difference Vegetation Index (NDVI)

3d picture : lat, long, index (ndvi)



Practicals with Ewen Gallic on wildfires and natural disasters

Convolutional Network

The light blue grid is called the input feature map. A kernel (shaded area) of value

$$\begin{pmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{pmatrix}$$

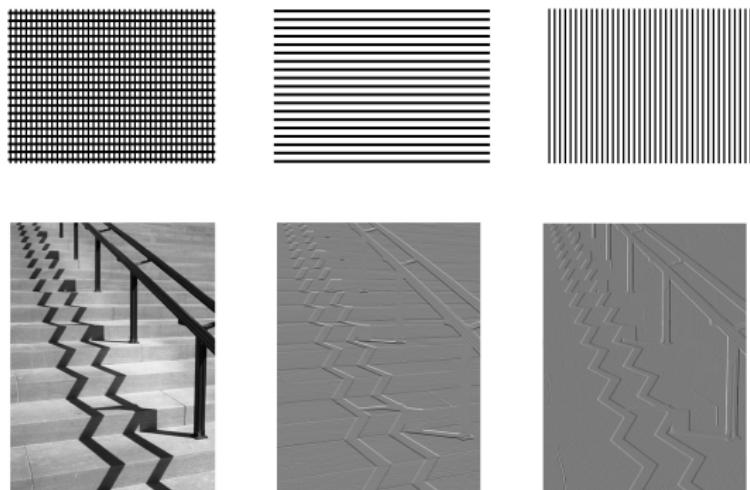
slides across the input feature map.
(see <http://deeplearning.net/>)

Extracting Shapes & Contours

See [Sobel Operator](#),

horizontal
$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

vertical
$$\begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$



Extracting Shapes & Contours

One can also compute gradients to extract contours



Picture Classifier

use <https://github.com/bnosac/image> to extract contours

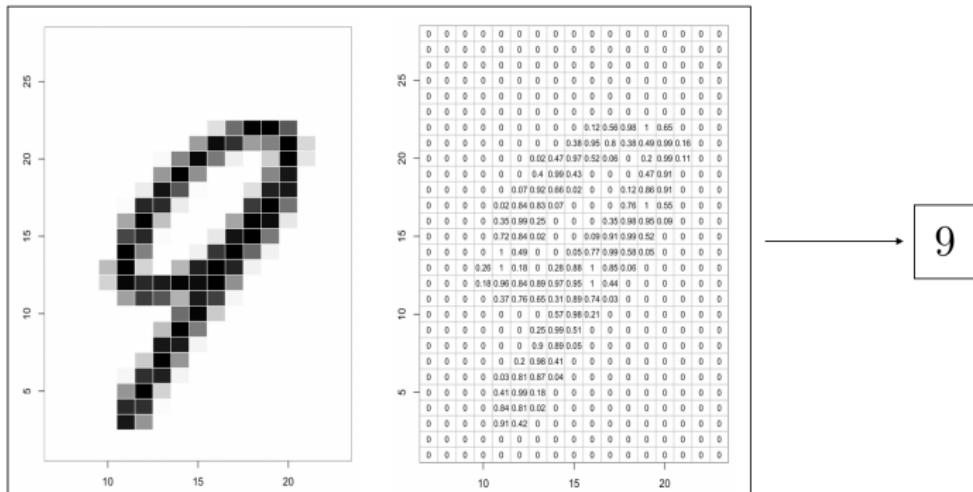
```
1 install.packages("image.LineSegmentDetector", repos =
  "https://bnosac.github.io/drat")
2 library(image.LineSegmentDetector)
3 library(imager)
4 img_path <- "../grey-lausanne.png"
5 im <- load.image(img_path)[,,1,1]
6 linesegments <- image_line_segment_detector(im*255)
7 plot(linesegments)
```

or

```
8 library(image.ContourDetector)
9 contourlines <- image_contour_detector(im*255)
10 plot(contourlines)
```

Picture Classifier

Classical $\{0, 1, \dots, 9\}$ handwritten digit recognition



$$x_i \in \mathcal{M}_{28,28}$$

$$y_i \in \{0, 1, \dots, 9\}$$

see tensorflow (Google) and keras.

Picture Classifier

See <https://www.tensorflow.org/install>

```
1 library(tensorflow)
2 install_tensorflow(method = "conda")
3 library(keras)
4 install_keras(method = "conda")

5 mnist <- dataset_mnist()
6 idx123 = which(mnist$train$y %in% c(1,2,3))
7 V <- mnist$train$x[idx123[1:800], ,]
8 MV = NULL
9 for(i in 1:800) MV=cbind(MV,as.vector(V[i,,]))
10 MV=t(MV)
11 df=data.frame(y=mnist$train$y[idx123][1:800],x=MV)
12 reg=glm((y==1)~.,data=df,family=binomial)
```

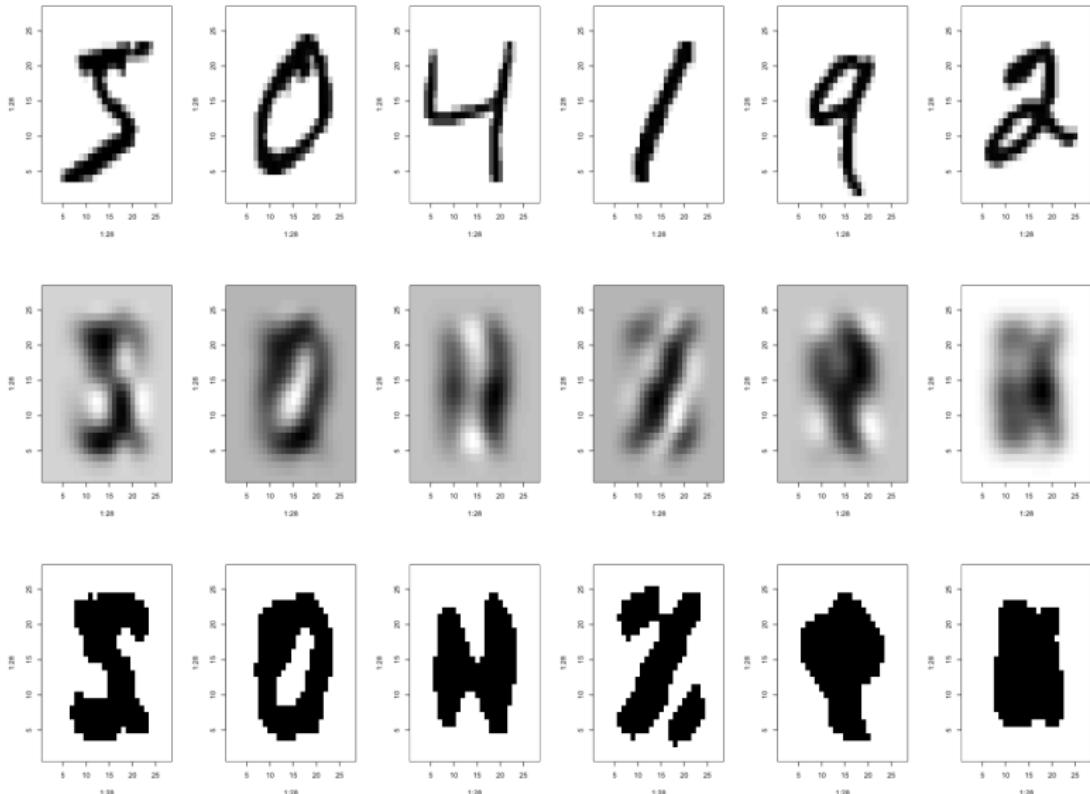
Here $x_i \in \mathbb{R}^{784}$ (for a small greyscale picture).

Picture Classifier

Still with those handwritten digit pictures, from the `mnist` dataset.
Here $\{(y_i, \mathbf{x}_i)\}$ with $y_i = "3"$ and $\mathbf{x}_i \in [0, 1]^{28 \times 28}$

One can define a [multilinear principal component analysis](#)
see also Tucker decomposition, from Hitchcock (1927, [The expression of a tensor or a polyadic as a sum of products](#))
We can use those decompositions to derive a simple classifier.

Picture Classifier



Picture Classifier

To reduce dimension use (classical) PCA

```
1 V <- (mnist$train$x[1:1000, ,])
2 MV <- NULL
3 for(i in 1:1000) MV <- cbind(MV, as.vector(V[i, ,]))
4 pca <- prcomp(t(MV))
```

or Multivariate PCA

```
5 library(rTensor)
6 T <- as.tensor(mnist$train$x[1:1000, ,])
7 tensor_svd <- hosvd(T)
8 tucker_decomp <- tucker(T, ranks = c(100, 3, 3))
9 T_approx <- ttl(tucker_decomp$Z, tucker_decomp$U, 1:3)
10 pca <- T_approx@data
```

or directly

```
11 pca <- mPCA(T)
```

Picture Classifier

with the (classical) PCA

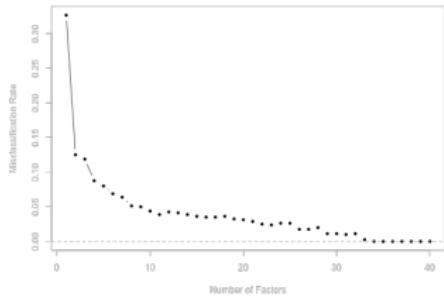
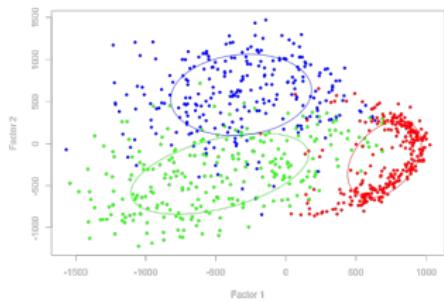
```
1 library(factoextra)
2 res.ind <- get_pca_ind(pca)
3 PTS <- res.ind$coord
4 value <- mnist$train$y[idx123][1:1000]
```

Consider a multinomial logistic regression, on the first k components

```
5 k <- 10
6 df <- data.frame(y=as.factor(value), x=PTS[,1:k])
7 library(nnet)
8 reg <- multinom(y~., data=df, trace=FALSE)
9 df$pred <- predict(reg, type="probs")
```

Picture Classifier

Consider a dataset with only three labels,
 $y_i \in \{1, 2, 3\}$ and the k principal components obtained by PCA,
 $\{x_1, \dots, x_{784}\} \rightarrow \{\tilde{x}_1, \dots, \tilde{x}_k\}$
E.g. scatterplot $\{\tilde{x}_{1,i}, \tilde{x}_{2,i}\}$
with ● when $y_i = 1$, ● $y_i = 2$ and ● $y_i = 3$,
Let m_k denote the multinomial logistic regression
on $\tilde{x}_1, \dots, \tilde{x}_k$ and visualize the misclassification rate



But one can also use some neural network model,
see Nielsen (2018, [Neural Networks and Deep Learning](#))

Consider optimization problem

$$\min_{\mathbf{x} \in \mathcal{X}} \{f(\mathbf{x})\}$$

solved by gradient descent,

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma_n \nabla f(\mathbf{x}_n), \quad n \geq 0$$

with starting point \mathbf{x}_0

γ_n can be call **learning rate**

Sometime the dimension of the dataset can be really big: we can't pass all the data to the computer at once to compute $f(\mathbf{x})$
we need to divide the data into smaller sizes and give it to our computer one by one

Batches

The Batch size is a hyperparameter that defines the number of sub-samples we use to compute quantities (e.g. the gradient).

Epoch

One Epoch is when the dataset is passed forward and backward through the neural network (only once).

We can divide the dataset of 2,000 examples into batches of 400 then it will take 5 iterations to complete 1 epoch.

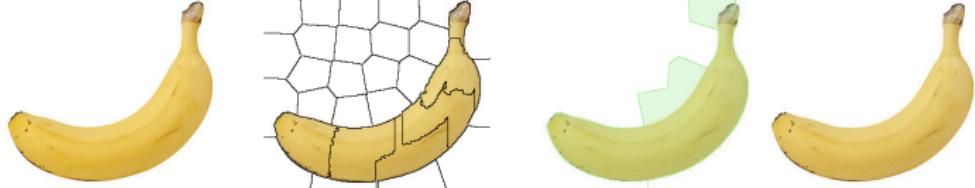
The number of epochs is the hyperparameter that defines the number times that the learning algorithm will work through the entire training dataset.

Picture Classifier

The algorithm is trained on tagged picture, i.e. (x_i, y_i) where here y_i is some class of fruits (banana, apple, kiwi, tomatoe, lime, cherry, pear, lychee, papaya, etc)



The challenge is to provide a new picture (x_{n+1} and to see the prediction \hat{y}_{n+1})



See [fruits-360/Training](#), e.g. the banana (118)

Picture Classifier

```
1 library(keras)
2 library(lime)
3 library(magick)
4 model <- application_vgg16(weights = "imagenet",
      include_top = TRUE)
5 test_image_files_path <- "/fruits-360/Test"
6 img <- image_read('https://upload.wikimedia.org/
      wikipedia/commons/thumb/8/8a/Banana-Single.jpg/272
      px-Banana-Single.jpg')
7 img_path <- file.path(test_image_files_path, "Banana",
      'banana.jpg')
8 image_write(img, img_path)
9 plot(as.raster(img))
10 plot_superpixels(img_path, n_superpixels = 35, weight
      = 10)
```

Our testing picture is from ://upload.wikimedia.org/

Picture Classifier

We use here (pre-trained) VGG16 and VGG19 models for Keras,
see Simonyan & Zisserman (2014, [Very Deep Convolutional
Networks for Large-Scale Image Recognition](#))

```
1 image_prep <- function(x) {  
2   arrays <- lapply(x, function(path) {  
3     img <- image_load(path, target_size = c(224,224))  
4     x <- image_to_array(img)  
5     x <- array_reshape(x, c(1, dim(x)))  
6     x <- imagenet_preprocess_input(x)  
7   })  
8   do.call(abind::abind, c(arrays, list(along = 1)))  
9 }  
10 res <- predict(model, image_prep(img_path))
```

To create our own model, see Chollet & Allaire (2018, [Deep
Learning with R](#)) [[github](#)]

Picture Classifier

```
1 imangenet_decode_predictions(res)
2 [[1]]
3   class_name  class_description      score
4 1  n07753592          banana 0.9929747581
5 2  n03532672          hook  0.0013420789
6 3  n07747607         orange 0.0010816196
7 4  n07749582         lemon 0.0010625814
8 5  n07716906 spaghetti_squash 0.0009176208
```

with 99.29% chance, the wikipedia picture of a banana is a *banana*

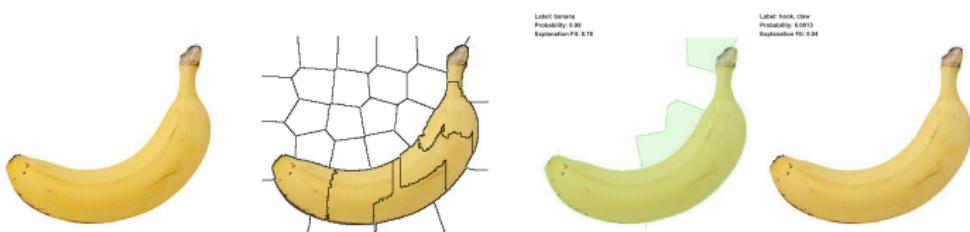


Picture Classifier

With the lime package - Local Interpretable Model-Agnostic Explanations we can also get some explanation of *why*

```
1 model_labels <- readRDS(system.file('extdata',  
2   imangenet_labels.rds', package = 'lime'))  
3 explainer <- lime(img_path, as_classifier(model,  
4   model_labels), image_prep)  
5 explanation <- explain(img_path, explainer,  
6   n_labels = 2, n_features = 35,  
7   n_superpixels = 35, weight = 10,  
8   background = "white")  
9 plot_image_explanation(explanation)
```

(see also the [vignette](#) or the [kitten](#) example)



Picture Classifier

We can build a **class activation map**, <http://gradcam.cloudcv.org> see Selvaraju *et al.* (2016, *Visual Explanations from Deep Networks via Gradient-based Localization*)

Compute the gradient of the score for class c , y_c (before softmax), with respect to feature maps A^k of convolutional layer, $\frac{\partial y_c}{\partial A^k}$

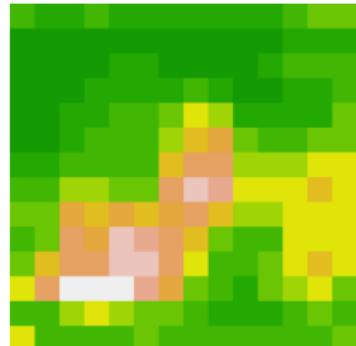
Define the importance weights as an average of gradients

$$\alpha_{c,k} = \frac{1}{n_x n_y} \sum_i \sum_j \frac{\partial y_c}{\partial A_{i,j}^k}$$

and define

$$\text{relu} \left(\sum_k \alpha_{c,k} A^k \right)$$

On a 14×14 grid, we obtain the picture on the right



Picture Classifier

```
1 img_path <- "car-accident.png"  
2 img <- image_load(img_path,  
3 target_size = c(224,224)) %>%  
4 image_to_array() %>%  
5 array_reshape(dim =  
6 c(1,224,224,3)) %>%  
7 imagenet_preprocess_input()
```

Training pictures are $224 \times 224 \times 3$ tensors
Convert to $224 \times 224 \times 3$ size



Picture Classifier

```
7 model <- application_vgg16(weights = "imagenet")  
8 preds <- model %>% predict(img)
```

Using the pre-trained **VGG16** network
on imangenet dataset (55,3467,096 pictures)

```
9 imagenet_decode_predictions(preds)  
10   class_name      class_desc      score  
11 1  n03594945      jeep          0.336  
12 2  n02814533      beach_wagon    0.197  
13 3  n03100240      convertible    0.153  
14 4  n03930630      pickup         0.062  
15 5  n03770679      minivan        0.041
```



Picture Classifier

```
2 img_path <- "house-fire.png"
```

Using the pre-trained **VGG16** network
on imangenet dataset (55,3467,096 pictures)

```
2 imangenet_decode_predictions(preds)
3   class_name  class_description score
4 1 n09472597      volcano    0.414
5 2 n09288635      geyser     0.320
6 3 n03773504      missile    0.104
7 4 n04310018 steam_locomotive 0.051
8 5 n04008634      projectile 0.045
```



Picture Classifier

```
2 img_path <- "ski-fall.jpg"
```

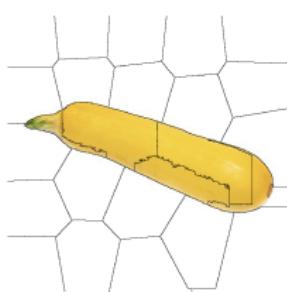
Using the pre-trained **VGG16** network
on imagenet dataset (55,3467,096 pictures)

```
2 imagenet_decode_predictions(preds)
3   class_name  class_description score
4 1 n04228054          ski 9.971e-01
5 2 n09193705          alp 2.449e-03
6 3 n04252077      snowmobile 8.625e-05
7 4 n02860847       bobsled 5.331e-05
8 5 n03218198      dogsled 4.959e-05
```

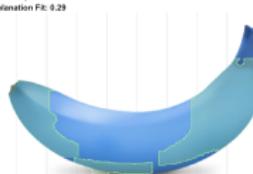
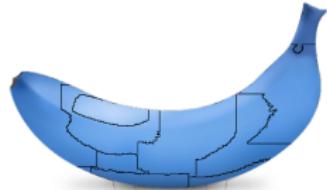


Picture Classifier

Note that (x_{n+1} can be something that was not in the training dataset (here yellow zucchini))



or some invented one (here a blue banana)



see also [pyimagesearch](#) for food/no food classifier

Picture Classifier

To train our own neural networks, use Chollet & Allaire (2018,
Deep Learning with R) [[github](#)]'s code (with `keras`)

```
1 channels <- 3
2 img_width <- 20
3 img_height <- 20
4 target_size <- c(img_width, img_height)
5 train_image_files_path <- "../fruits-360/Training/"
6 valid_image_files_path <- "../fruits-360/Validation/"
7 library(keras)
8 library(lime)
9 library(magick)
10 library(ggplot2)
11 train_samples <- train_image_array_gen$n
12 valid_samples <- valid_image_array_gen$n
```

Picture Classifier

```
1 train_data_gen = image_data_generator(  
2     rescale = 1/255  
3     # rotation_range = 40,  
4     # width_shift_range = 0.2,  
5     # height_shift_range = 0.2,  
6     # shear_range = 0.2,  
7     # zoom_range = 0.2,  
8     # horizontal_flip = TRUE  
9 )  
10 batch_size <- 32  
11 epochs <- 10
```

Picture Classifier

```
1 train_image_array_gen <- flow_images_from_directory(  
2     train_image_files_path,  
3     train_data_gen,  
4     target_size = target_size,  
5     class_mode = "categorical",  
6     classes = fruit_list)  
7 valid_data_gen <- image_data_generator(rescale =  
8     1/255)  
9 valid_image_array_gen <- flow_images_from_directory(  
10    valid_image_files_path,  
11    valid_data_gen,  
12    target_size = target_size,  
13    class_mode = "categorical",  
14    classes = fruit_list)
```

Picture Classifier

```
1 model <- keras_model_sequential()
2 # add layers
3 model %>%
4   layer_conv_2d(filter = 32, kernel_size = c(3,3),
5     padding = "same", input_shape = c(img_width,
6       img_height, channels)) %>%
7   layer_activation("relu") %>%
8
9   # Second hidden layer
10  layer_conv_2d(filter = 16, kernel_size = c(3,3),
11    padding = "same") %>%
12    layer_activation_leaky_relu(0.5) %>%
13    layer_batch_normalization() %>%
14
15    # Use max pooling
16    layer_max_pooling_2d(pool_size = c(2,2)) %>%
17    layer_dropout(0.25) %>%
```

Picture Classifier

```
15 # Flatten max filtered output into feature vector
16 # and feed into dense layer
17 layer_flatten() %>%
18 layer_dense(100) %>%
19 layer_activation("relu") %>%
20 layer_dropout(0.5) %>%
21
22 # Outputs from dense layer are projected onto output
23 # layer
24 layer_dense(output_n) %>%
25 layer_activation("softmax")
26
27 # compile
28 model %>% compile(
29   loss = "categorical_crossentropy",
30   optimizer = optimizer_rmsprop(lr = 0.0001, decay = 1
31     e-6),
32   metrics = "accuracy"
33 )
```

Picture Classifier

```
1 hist <- model %>% fit_generator(  
2   # training data  
3   train_image_array_gen,  
4   # epochs  
5   steps_per_epoch = as.integer(train_samples /  
6     batch_size),  
7   epochs = epochs,  
8   # validation data  
9   validation_data = valid_image_array_gen,  
10  validation_steps = as.integer(valid_samples /  
11    batch_size),  
12  # print progress  
13  verbose = 2,  
14  callbacks = list(  
15    # save best model after every epoch  
16    callback_model_checkpoint("fruits_checkpoints.h5",  
17      save_best_only = TRUE),  
18    # only needed for visualising with TensorBoard  
19    callback_tensorboard(log_dir = "logs")))
```

Facial Recognition

For facial recognition, see Schwemmer (2018, `facerec: An interface for face recognition in R`), based on `kairos`



or <https://github.com/bnosac/image> for various R functions
(Corner Dor Edges etection, etc).

Picture Classifier

```
1 install.packages("image.darknet", repos = "https://
    bnosac.github.io/drat")
2 library(image.darknet)
3 yolo_tiny_voc <- image_darknet_model(
4     type = 'detect',
5     model = "tiny-yolo-voc.cfg",
6     weights = system.file(
7         package = "image.darknet", "models", "tiny-yolo-voc.
    weights"),
8     labels = c("aeroplane", "bicycle", "bird", "boat", " "
    bottle", "bus", "car", "cat", "chair", "cow", " "
    diningtable", "dog", "horse", "motorbike", "person
    ", "pottedplant", "sheep", "sofa", "train", " "
    tvmonitor"))
9 image_darknet_detect(file = "lausanne-d-1.png",
10     object = yolo_tiny_voc)
```

Picture Classifier

```
11 Loading weights from /Library/Frameworks/R.framework/
   Versions/3.5/Resources/library/image.darknet/
   models/tiny-yolo-voc.weights...Done!
12 lausanne-d-1.png: Predicted in 0.904149 seconds.
13 Boxes: 845 of which 6 above the threshold.
14 person: 82%
15 person: 47%
16 person: 61%
17 person: 60%
18 person: 54%
19 person: 40%
```